

CENG 795

Advanced Ray Tracing

Fall '2022-2023

Assignment 5 - Advanced Lighting and HDR Rendering
(v.1.0)

Due date: January 2, 2022, Monday, 23:59



1 Objectives

In this homework you are going to add new lights to your ray tracers. Some of these lights can be high dynamic range (HDR) environment maps. As a result of this, you will no longer save your outputs as PNG files clipped to $[0, 255]$ range. Instead you will save your results as HDR images in .exr or .hdr formats. You will also tone map these images and then save them as PNG images. The following new lights are added:

- Directional lights
- Spot lights
- Environment lights

Keywords: *environment mapping, HDR imaging, gamma correction*

2 Specifications

The specifications are the same as for the previous homeworks so are not repeated here. For reading/writing HDR images you can use a library of your choice. OpenCV has support for HDR images. There appears to be another library (not tested) called `tinyexr` that you can use for this purpose: <https://github.com/syoyo/tinyexr>

3 Scene File

Please see the previous assignments for a detailed description of the scene file. In this section, only the elements introduced in this homework are explained.

3.1 Camera

Camera element has now the following added child elements to support tone mapping:

```
<ImageName>output.exr</ImageName>
<Tonemap>
  <TMO>Photographic</TMO>
  <TMOOptions>0.18 1</TMOOptions> <!-- key_value burn_percent -->
  <Saturation>1.0</Saturation>
  <Gamma>2.2</Gamma>
</Tonemap>
```

`ImageName` element existed before but now it can contain an HDR file name. If this is the case (as in `output.exr`), you must use the following tone mapping specifications. You are only expected to support the photographic tone mapping algorithm¹. After tone mapping, you should save a second image by replacing the `.exr` extension with `.png`. So in this case, you should save both `output.exr` and `output.png`.

Tone mapping operators generally compress the luminance of the HDR image, leaving the colors untouched. You can use the following pipeline for tone mapping:

1. Compute luminance from RGB: $Y_i = 0.2126R + 0.7152G + 0.0722B$. There can be other conversions but this conversion assumes that the input RGB colors are in sRGB space, which is a reasonable assumption.
2. Apply the tone mapping algorithm: $Y_o = f(Y_i)$, where f represents the tone mapping function. The details of this process are in lecture slides and in the paper. You can use the global version of the operator for simplicity.
3. For the previous step, use the key parameter given as the first value of the `TMOOptions` element. Also, if the burn-out percentage parameter (the second value of `TMOOptions`) is not zero, it means you must apply some burn-out. For this purpose, you must sort the input

¹Reinhard, E., Stark, M., Shirley, P., & Ferwerda, J. (2002, July). Photographic tone reproduction for digital images. In Proceedings of the 29th annual conference on Computer graphics and interactive techniques (pp. 267-276). <https://web.tecgraf.puc-rio.br/~scuri/inf1378/pub/reinhard.pdf>

luminances (e.g., `std::sort`) and find the luminance value corresponding to the given burn-out percentage. For example, if this value is 1, it means you must select the 99th percentile luminance value as the burn-out threshold. This is the L_{white} value to be used in Equation 4 of the original paper. If burn-out percentage is zero, you can simply use Equation 3 of the paper. In either case, do not forget that you must first apply Equations 1 and 2.

4. Next, you must compute the new RGB colors from the compressed luminance value, Y' . This process can be formulized as follows:

$$R_o = Y_o \left(\frac{R}{Y_i} \right)^s \quad G_o = Y_o \left(\frac{G}{Y_i} \right)^s \quad B_o = Y_o \left(\frac{B}{Y_i} \right)^s. \quad (1)$$

Here s represents the saturation parameter value from the XML file. The closer this value to 0, the more grayish the output image will look like. Usually, the value of this parameter will be 1. If you increase it you can obtain more saturated results.

5. Finally, you must perform gamma correction for storing this image as a low dynamic range PNG file. For this purpose first clamp all of the RGB values that you found above to the $[0, 1]$ range. Then apply the following equations, where g represent the gamma parameter in the XML file:

$$R_f = 255 (R_o)^{(1/g)} \quad G_f = 255 (G_o)^{(1/g)} \quad B_f = 255 (B_o)^{(1/g)}. \quad (2)$$

You can then round these values to the nearest integer and store them in the PNG file as you have done in previous homeworks.

3.2 Lights

The lights element may now contain directional lights, spot lights, and spherical directional lights (i.e., environment maps) in addition to the previous lights. Directional lights are simply defined by their direction and radiance as shown below. As they are defined by their radiance, there is no distance based attenuation for these lights.

```
<DirectionalLight id="1">
  <Direction>1 -0.8 -1</Direction>
  <Radiance>200 200 200</Radiance>
</DirectionalLight>
```

Spot lights are similar to point lights but they send all their energy within a predefined cone of illumination. This cone is defined by two angles, the coverage angle and the fall-off angle. Within the fall-off angle, the light behaves exactly as a point light. This is shown as the no-fall-off zone in Figure 1. Between the fall-off and coverage angles, the light intensity will diminish non-linearly. Beyond the coverage angle, the light does not provide any illumination. A sample definition is shown below:

```
<SpotLight id="1">
  <Position>-0.93 1 0.9</Position>
  <Direction>1 -1 -1</Direction>
  <Intensity>600 600 600</Intensity>
```




Figure 1: Spot light geometry.

```
<CoverageAngle>10</CoverageAngle>
<FalloffAngle>8</FalloffAngle>
</SpotLight>
```

This light is attenuated based on its distance to the point being shaded just like a regular point light. For spot attenuation, you can use the following formula:

$$s = \left(\frac{\cos(\alpha) - \cos(\text{falloff}/2)}{\cos(\text{falloff}/2) - \cos(\text{coverage}/2)} \right)^4, \quad (3)$$

where α is the angle between the light direction and the point being shaded. The net irradiance at a point x that is d units away from a spot light is then given by:

$$E(x) = \begin{cases} \frac{I}{d^2} & \text{if } x \text{ is within the no-fall-off zone,} \\ s \frac{I}{d^2} & \text{if } x \text{ is outside the no-fall-off zone but inside the coverage zone,} \\ 0 & \text{otherwise.} \end{cases} \quad (4)$$

The last type of light that you need to support is the spherical directional light. This light provides illumination from all angles: it is an environment light. The radiance values for this light are defined by an HDR image in latitude-longitude format. Therefore, the definition of this light uses both a texture element and a light element. An example is shown below:

```
<Textures>
  <Images>
    <Image id="1">textures/pisa_latlong.exr</Image>
  </Images>
</Textures>
<Lights>
  <SphericalDirectionalLight id="1">
    <ImageId>1</ImageId>
  </SphericalDirectionalLight>
</Lights>
```

To support this light, you need to sample a random direction $d = (x, y, z)$ from the upper hemisphere of the surface. You can use random rejection sampling for this purpose (do not worry if your results are a bit noisier than the provided images at this point). Once a direction is chosen (make




Figure 2: The same car model is illuminated by different environment maps giving rise to very different appearances.

sure that it is a unit vector), you need to convert it to (u, v) values and look up the corresponding value in the HDR environment map. You can use the following formula for this purpose:

$$u = \frac{1 + \frac{\text{atan2}(d_x, d_z)}{\pi}}{2}, \quad (5)$$

$$v = \frac{\text{acos}(d_y)}{\pi}, \quad (6)$$

$$L(d) = I(u, v), \quad (7)$$

where I represent the HDR environment map. As with area lights, you need to compensate for the fact that you sampled a single direction for illumination. This is done by dividing the fetched radiance value by the probability density of selecting that particular direction. Because we sample a random direction uniformly, this probability density is equal to $p(d) = 1/(2\pi)$. Thus the overall radiance estimate is computed by:

$$\hat{L} = \frac{L(d)}{1/(2\pi)} = 2\pi L(d) \quad (8)$$

Note that some spectacular results can be obtained by using this technique. Figure 2 shows two renderings of the same scene with different environment maps. The appearance of the car model changes entirely as if it is physically present in different scenes. Achieving such a result with point lights could be very difficult, if not impossible.

4 Hints & Tips

In addition to those for the previous homeworks, the following tips may be useful for this homework.

1. There are a few scenes among the input files that will make you grateful for having implemented an acceleration structure, but regretful that you haven't optimized it further.
2. The car model has a few degenerate triangles in some PLY files. I used to clean such triangles in the past, but it is important that you can ignore them in your own ray tracers as advanced scenes generally have such cases. When you insert a triangle to a mesh, you can simply ignore that triangle if at least any two vertices of the triangle have the same position.

3. As the car model has some intentional cracks between different parts of the car, it is a good idea to render this model without backface culling as otherwise, you will see through the car between these cracks.

5 Bonus

I will be more than happy to give bonus points to students who make important contributions such as new scenes, importers/exporters between our XML format and other standard file formats. Note that a Blender exporter², which exports Blender data to our XML format, was written by one of our previous students. You can use this for designing a scene in Blender and exporting it to our file format.

6 Regulations

1. **Programming Language:** C/C++ is the recommended language. However, other languages can be used if so desired. In the past, some students used Rust or even Haskell for implementing their ray tracers.
2. **Changing the Sample Codes:** You are free to modify any sample code provided with this homework.
3. **Additional Libraries:** If you are planning to use any library other than *(i)* the standard library of the language, *(ii)* pthread, *(iii)* the XML parser, and the PNG libraries please first ask about it on ODTUClass and get a confirmation. Common sense rules apply: if a library implements a ray tracing concept that you should be implementing yourself, do not use it!
4. **Submission:** Submission will be done via ODTUClass. To submit, Create a “**tar.gz**” file named “raytracer.tar.gz” that contains all your source code files and a Makefile. The executable should be named as “raytracer” and should be able to be run using the following commands (scene.xml will be provided by us during grading):

```
tar -xf raytracer.tar.gz
make
./raytracer scene.xml
```

Any error in these steps will cause point penalty during grading.

5. **Late Submission:** You can submit your codes up to 3 days late. Each late day will cause a 10 point penalty.
6. **Cheating: We have zero tolerance policy for cheating.** People involved in cheating will be punished according to the university regulations and will get 0 from the homework. You can discuss algorithmic choices, but sharing code between groups or using third party code is strictly forbidden. By the nature of this class, many past students make their ray tracers publicly available. You must refrain from using them at all costs.

²<https://saksagan.ceng.metu.edu.tr/courses/ceng477/student/ceng477exporter.py>

7. **Forum:** Check the ODTUClass forum regularly for updates/discussions.
8. **Evaluation:** The basis of evaluation is your blog posts. Please try to create interesting and informative blog posts about your ray tracing adventures. You can check out various past blogs for inspiration. However, also expect your codes to be compiled and tested on some examples for verification purposes. So the images that you share in your blog post must directly correspond to your ray tracer outputs.