

CENG 795

Advanced Ray Tracing

Fall '2022-2023

Assignment 6 - Putting it All Together: BRDFs and Path Tracing
(v.1.0)

Due date: January 22, 2022, Sunday, 23:59



1 Objectives

In this homework we are going to bring together all the things that we learned so far to implement a fully functional path tracer. In particular the features that you will implement are:

- BRDF models
- Object lights
- Path tracing with uniform and important sampling

Keywords: *path tracing, BRDF, importance sampling*

2 Specifications

The specifications are the same as for the previous homeworks so are not repeated here. Make sure you check out the supporting PDF files that are separately uploaded in case you need extra clarifications.

3 Scene File

3.1 BRDF

A BRDF will be defined in the scene file as follows:

```
<BRDFs>
  <ModifiedBlinnPhong id="1" normalized="true">
    <Exponent>50</Exponent>
  </ModifiedBlinnPhong>
</BRDFs>
```

Possible BRDF types are:

- OriginalBlinnPhong (this is the BRDF of the default shading model we have been using so far)
- OriginalPhong
- ModifiedBlinnPhong (can be normalized as shown above)
- ModifiedPhong (can be normalized)
- TorranceSparrow (normalized by definition)

For the Torrance-Sparrow BRDF the exponent is the p term used in the distribution function. For that BRDF, there is an attribute called *kdfresnel*. If this is true, it means that you should use $\frac{(1-F)k_d}{\pi}$ instead of $\frac{k_d}{\pi}$ when computing the diffusely reflected radiance. Please refer to the BRDF document for specific details.

3.2 Object Lights

Any regular object can now be a light source. For example, we can make a sphere light source by defining it as **LightSphere**:

```
<LightSphere id="1">
  <Material>1</Material>
  <Center>11</Center>
  <Radius>0.2</Radius>
  <Radiance>31.831 31.831 31.831</Radiance>
</LightSphere>
```

This is the same definition as a regular sphere except that it also has a radiance element. A mesh light is similar:

```

<LightMesh id="1">
  <Faces> 2 6 7 7 3 2 </Faces>
  <Material>1</Material>
  <Radiance>1.21 1.21 1.21</Radiance>
</LightMesh>

```

Note that in case of pure path tracing, you do not have to do anything special to sample an object light. Some rays will simply hit them and get the corresponding radiance value (still divided by the probability of uniform or importance sampling). But if you are doing explicit light sampling (also known as *next event estimation*), you must sample a point on the light source and divide the returned radiance value by the probability of sampling the corresponding direction.

3.3 Path Tracing

Until now, we sent a recursive ray only when an object was a mirror, conductor, or dielectric. With path tracing, we send recursive rays for all types of objects. This is because in reality all objects exchange light with each other. We may call these recursively sent rays as global or indirect illumination rays. This way if only a part of the scene receives direct illumination, the scene will still be illuminated because that part will reflect light to other parts.

With path tracing, we typically disable the crude ambient term we used so far. Instead, each scene point receives some ambient light depending on its geometric relationship to other objects in the scene. In the teaser figure shown at the top, almost everywhere would be black without path tracing because there is almost no direct illumination: all the light comes from a small opening behind the door (hence the name *ajar*).

Path tracing mode is enabled through the addition of the following XML elements to the **Camera** element:

```

<Renderer>PathTracing</Renderer>
<RendererParams>NextEventEstimation ImportanceSampling RussianRoulette</RendererParams>

```

So we can define multiple cameras, each camera with a different renderer setting. The **RendererParams** element is optional. If it is missing or empty, it means path tracing with uniform sampling. Otherwise it may contain the following options:

- Importance Sampling: you must use cosine based sampling of the upper hemisphere instead of uniform sampling.
- Next Event Estimation: you must directly sample light sources in addition to random sampling (see below)
- Russian Roulette: you must terminate ray paths probabilistically rather than using a fixed recursion depth (see below)

In next event estimation, for each primary ray we perform two operations: (1) Sample a direction that will hit a light source (in case of an occluder the hit may not happen, but this is not relevant) and (2) sample a random global illumination ray. If the global illumination ray happens to hit the same light source that you explicitly sampled in (1), you must discard its contribution. Otherwise, you count a light twice and introduce bias in your result.

Russian roulette is a dangerous game for humans and rays alike. The idea is to probabilistically terminate a ray path if it is likely to make a low contribution. For example, after many bounces

the impact of the ray will significantly diminish because each bounce will absorb the energy of the ray. Assume that initially the ray has a *throughput* of 1. If a ray bounces from a surface along a direction whose BRDF is 0.2, its throughput will be 0.2. Now if the ray bounces again from another surface with a BRDF of 0.1, the new throughput will be 0.02. This means, whatever radiance that may come from the direction of the ray from this point on, it will get scaled by a factor of 0.02. So we can use this throughput factor to probabilistically terminate the ray. The steps to implement Russian roulette are as follows:

1. Draw a normalized uniform number, ξ
2. If $\xi > q$ where $1 - q$ is the termination probability, kill the rest of this path (e.g., if $q = 0.9$, the ray will be terminated with 0.1 probability). You can use any q you want but using throughput is a sensible choice: low throughput paths are more likely to be killed than high throughput paths.
3. Add a constant term C to model the radiance that we may have lost due to terminating the path. $C = 0$ is a common choice.
4. Compute the final estimate as a weighted average of the computed estimate and the ignored estimate:

$$F' = \begin{cases} \frac{F - (1-q)C}{q}, & \xi \leq q \\ C, & \xi > q \end{cases} \quad (1)$$

Note that the expected value of F' is the same as the expected value of F .

Note that variance will generally increase with Russian roulette - so it is not a variance reduction technique. But it gives us a statistically sound way to terminate rays instead of cutting them after a fixed number of bounces. In general, it is best to combine it with the maximum recursion depth parameter: kill the ray if it fails the throughput test *and* it exceeds the maximum recursion depth. However, the scene results shared with you do not implement this technique - only the direct application of Russian roulette is shown, therefore results are generally noisier.

4 Hints & Tips

In addition to those for the previous homeworks, the following tips may be useful for this homework.

1. There are a few scenes among the input files that will make you grateful for having implemented an acceleration structure, but regretful that you haven't optimized it further.
2. You can stratify your global illumination ray samples for reducing variance.
3. If you want, you can send multiple global illumination ray samples for each primary ray sample. But the shared results are created by sending one illumination sample per primary ray sample.
4. The sponza scene contains vertex and texture data in binary files. For vertex data it is just x, y, z values of each vertex packed together. You must open these files in binary mode to read them correctly. For texture data, it is the u, v values packed in the same manner. For

index data, it is the integer index values. Each binary data starts with a single integer value that represents the number of elements to read. For vertex data, it is the number of vertices. For texture data, it is the number of texture coordinates. For index data, it is the number of faces (in this case you must read three times of this value as each face is represented using 3 indices).

5. If you come across `<ZeroBasedIndexing>true</ZeroBasedIndexing>` element in the XML file, you must assume that indices are zero-based as opposed to the one-based convention that we have been using.
6. Similarly, if you come across `vertexOffset` or `textureOffset` attributes in face definitions, it means you must add these amounts to the indices before accessing the vertex and texture data. Note that these numbers can be negative!

5 Bonus

I will give a significant bonus to students who convert some complex PBRT (<https://github.com/mmp/pbrt-v4-scenes>) or Tungsten (<https://benedikt-bitterli.me/tungsten.html>) scenes to our format and render them correctly.

6 Regulations

1. **Programming Language:** C/C++ is the recommended language. However, other languages can be used if so desired. In the past, some students used Rust or even Haskell for implementing their ray tracers.
2. **Changing the Sample Codes:** You are free to modify any sample code provided with this homework.
3. **Additional Libraries:** If you are planning to use any library other than *(i)* the standard library of the language, *(ii)* pthread, *(iii)* the XML parser, and the PNG libraries please first ask about it on ODTUClass and get a confirmation. Common sense rules apply: if a library implements a ray tracing concept that you should be implementing yourself, do not use it!
4. **Submission:** Submission will be done via ODTUClass. To submit, Create a “**tar.gz**” file named “raytracer.tar.gz” that contains all your source code files and a Makefile. The executable should be named as “raytracer” and should be able to be run using the following commands (scene.xml will be provided by us during grading):

```
tar -xf raytracer.tar.gz
make
./raytracer scene.xml
```

Any error in these steps will cause point penalty during grading.

5. **Late Submission:** You can submit your codes up to 3 days late. Each late day will cause a 10 point penalty.

6. **Cheating: We have zero tolerance policy for cheating.** People involved in cheating will be punished according to the university regulations and will get 0 from the homework. You can discuss algorithmic choices, but sharing code between groups or using third party code is strictly forbidden. By the nature of this class, many past students make their ray tracers publicly available. You must refrain from using them at all costs.
7. **Forum:** Check the ODTUClass forum regularly for updates/discussions.
8. **Evaluation:** The basis of evaluation is your blog posts. Please try to create interesting and informative blog posts about your ray tracing adventures. You can check out various past blogs for inspiration. However, also expect your codes to be compiled and tested on some examples for verification purposes. So the images that you share in your blog post must directly correspond to your ray tracer outputs.